

RAG Document Processing Test

Modern document processing systems combine text extraction, image understanding, optical character recognition, and semantic retrieval to transform heterogeneous documents into useful knowledge. A typical pipeline first parses the source document, identifies embedded visual content, applies OCR where required, normalizes the extracted information, and finally prepares the content for chunking and retrieval. This test document intentionally contains normal text, regular images, screenshots containing text, and a simple processing diagram so that each stage of a document intelligence or RAG ingestion pipeline can be validated.

1. Normal PNG/JPG Images

Image 1 — Document Processing

A document ingestion pipeline prepares files for downstream analysis.

Image 2 — OCR

Optical character recognition extracts text from scanned pages and images.

Image 3 — Semantic Search

Embeddings represent document chunks as vectors for similarity search.

Image 4 — Retrieval

A retriever selects relevant chunks before the generation stage.

2. Screenshots Containing Text

OCR Processing Log

```
INFO Document uploaded successfully
INFO Rendering page 1 as image
INFO OCR started: language=eng
INFO Text extracted: 428 characters
INFO OCR completed successfully
```

Application Console

```
POST /api/documents/upload      200 OK
Parsed documentId=101
Detected images=4
Detected text screenshots=2
Processing pipeline completed
```

3. Optional Diagram



The diagram represents a simplified flow from document ingestion through parsing, OCR, chunking, and retrieval.